

# Computational Intelligence Laboratory

## Lecture 3

### Matrix Decomposition and Optimization

**Thomas Hofmann**

ETH Zurich – `cil.inf.ethz.ch`

Credit for slides and figures on optimization:

Martin Jaggi @EPFL

March 17, 2017

# Section 1

## Matrix Decomposition – Reloaded

# Beyond Singular Value Decomposition

- ▶ Is SVD the final answer for (low-rank) matrix decomposition?
- ▶ **Eckart-Young theorem** guarantees:

$$\mathbf{A}_k = \arg \min_{\text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_F^2$$

- ▶ surprising: not a convex optimization problem!
- ▶ convex combination of  $k$ -rank matrices is generally not rank  $k$

$$\underbrace{\frac{1}{2} \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}}_{\text{rank 1}} + \underbrace{\frac{1}{2} \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}}_{\text{rank 1}} = \underbrace{\frac{1}{2} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}}_{\text{rank 2}}$$

# Beyond Singular Value Decomposition

- ▶ Problem: entries which are **unobserved** – not zero!
  - ▶ should optimize

$$\min_{\text{rank}(\mathbf{B})=k} \left[ \sum_{(i,j) \in \mathcal{I}} (a_{ij} - b_{ij})^2 \right], \quad \mathcal{I} = \{(i, j) : \text{observed}\}$$

- ▶ instead of

$$\min_{\text{rank}(\mathbf{B})=k} \left[ \sum_{i,j} (a_{ij} - b_{ij})^2 \right] = \min_{\text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_F^2$$

- ▶ usually: mean zero  $a_{ij} \leftarrow a_{ij} - \frac{1}{|\mathcal{I}|} \sum_{\mathcal{I}} a_{ij}$

# Matrix Factorization: Non-Convex Problem

- ▶ Singular Value Decomposition is not enough!
- ▶ **Non-convex** optimization problem
  - ▶ variant A: non-convex constraint set

minimize over set  $\{\mathbf{B} : \text{rank}(\mathbf{B}) = k\}$

- ▶ variant B: non-convex objective

re-parametrize  $\mathbf{B} = \mathbf{UV}$ ,  $\mathbf{U} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{V} \in \mathbb{R}^{k \times n}$

then  $\text{rank}(\mathbf{B}) \leq k$

e.g.  $f(u, v) = (a - uv)^2$ ,  $u_1 v_1 = u_2 v_2 = a \wedge u_1 \neq u_2$

$$\implies f(u_1, v_1) = f(u_2, v_2) = 0 \wedge f\left(\frac{u_1 + u_2}{2}, \frac{v_1 + v_2}{2}\right) > 0$$

# Alternating Minimization

- ▶ Is there something **convex** about the **non-convex** objective?

$$f(\mathbf{U}, \mathbf{V}) = \frac{1}{|\mathcal{I}|} \sum_{(i,j) \in \mathcal{I}} (a_{ij} - \langle \mathbf{u}_i, \mathbf{v}_j \rangle)^2$$

- ▶ for fixed  $\mathbf{U}$ ,  $f$  is convex in  $\mathbf{V}$  – for fixed  $\mathbf{V}$ ,  $f$  is convex in  $\mathbf{U}$
- ▶ ... which does not mean  $f$  is jointly convex in  $\mathbf{U}$  and  $\mathbf{V}$
- ▶ Idea: perform **alternating minimization**

$$\mathbf{U} \leftarrow \arg \min_{\mathbf{U}} f(\mathbf{U}, \mathbf{V})$$

$$\mathbf{V} \leftarrow \arg \min_{\mathbf{V}} f(\mathbf{U}, \mathbf{V}), \quad \text{repeat until convergence}$$

- ▶  $f$  is never increased and lower bounded by 0

# Alternating Least Squares

- ▶ Alternating minimization for low-rank matrix factorization = **alternating least squares**
  - ▶ decompose  $f$  into subproblems for columns of  $\mathbf{V}$

$$f(\mathbf{U}, \mathbf{V}) = \sum_i \underbrace{\left[ \sum_{j:(i,j) \in \mathcal{I}} (a_{ij} - \langle \mathbf{u}_j, \mathbf{v}_i \rangle)^2 \right]}_{=: f(\mathbf{U}, \mathbf{v}_i)}$$

- ▶ least squares problem  $f(\mathbf{U}, \mathbf{v}_i)$  for column  $\mathbf{v}_i$  of  $\mathbf{V}$ 
  - ▶ each of which can be solved independently!
- ▶ by symmetry: also holds for  $\mathbf{U} \leftrightarrow \mathbf{V}$

# Frobenius Norm Regularization

- ▶ Typically: regularize matrix factors  $\mathbf{U}, \mathbf{V}$
- ▶ Frobenius norm regularizer

$$\Omega(\mathbf{U}, \mathbf{V}) = \|\mathbf{U}\|_F + \|\mathbf{V}\|_F$$

- ▶ then

$$\text{minimize} \rightarrow f(\mathbf{U}, \mathbf{V}) + \mu \Omega(\mathbf{U}, \mathbf{V}), \quad \mu > 0$$

- ▶ does not change separability structure of problem



# ALS for Collaborative Filtering

- ▶ given low-dimensional representations for items
- ▶ compute for each user independently the best representation
- ▶ given low-dimensional representations for users
- ▶ compute for each item independently the best representation
- ▶ all optimization problems are small least-square problems

# Matrix Decomposition as Optimization

- ▶ Is this the best we can do?
- ▶ Does this strategy always work (more factorizations to come ...)?
- ▶ We need to better understand the power of **convex optimization**!

## Section 2

# Unconstrained Optimization

# Optimization

- ▶ General optimization problem (**unconstrained minimization**)

$$\begin{array}{ll}\text{minimize} & f(\mathbf{x}) \\ \text{with} & \mathbf{x} \in \mathbb{R}^m\end{array}$$

- ▶ solutions  $\mathbf{x} \in \mathbb{R}^m$  (e.g. parameters in learning)
- ▶ objective  $f : \mathbb{R}^m \rightarrow \mathbb{R}$
- ▶ technical assumption:  $f$  is continuous and differentiable

# Why? And How?

Optimization is everywhere: *machine learning, big data, statistics, data analysis of all kinds, finance, logistics, planning, control theory, mathematics, search engines, simulations, and many other applications ...*


- ▶ **Mathematical Modeling:**
  - ▶ defining & modeling learning as optimization problems
- ▶ **Computational Optimization:**
  - ▶ designing & running an (appropriate) optimization algorithm

# Optimization Algorithms

- ▶ Optimization at large scale: **simplicity** rules!
- ▶ Main approaches:
  - ▶ Coordinate Descent
  - ▶ Gradient Descent
  - ▶ Stochastic Gradient Descent (SGD)
- ▶ History:
  - ▶ 1950s: Linear Programming
  - ▶ 1980s: General optimization, convex optimization theory
  - ▶ 2005-today: Large scale optimization

# Coordinate Descent

**Goal:** Find  $\mathbf{x}^* = \arg \min_{\mathbf{x}} f(\mathbf{x})$ .



**Idea:** update one coordinate at a time, keeping all others fixed.

# Coordinate Descent

initialize  $\mathbf{x}^{(0)} \in \mathbb{R}^m$

**for**  $t = 0, \dots, T-1$  **do**

sample coordinate  $d \stackrel{\text{uni}}{\sim} \{1 \dots m\}$

**optimize** (analytically or via line search)

$$u^* \leftarrow \arg \min_{u \in \mathbb{R}} f(x_1^{(t)}, \dots, x_{d-1}^{(t)}, \textcolor{red}{u}, x_{d+1}^{(t)}, \dots, x_m^{(t)})$$

**update**

$$x_i^{(t+1)} \leftarrow \begin{cases} u^* & \text{if } i = \textcolor{red}{d} \\ x_i^{(t)} & \text{otherwise} \end{cases}$$


**end for**



# Gradient

**Gradient** of a function  $f : \Omega \subseteq \mathbb{R}^m \rightarrow \mathbb{R}$ :

$$\nabla f := \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \dots \\ \frac{\partial f}{\partial x_m} \end{pmatrix}, \quad \nabla f : \Omega \rightarrow \mathbb{R}^m$$



# Steepest Descent


- ▶ first suggested by Cauchy in 1847
- ▶ simple to implement, scalable and robust

initialize  $\mathbf{x}^{(0)} \in \mathbb{R}^m$

**for**  $t = 0:T-1$  **do**

    update  $\mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \eta \nabla f(\mathbf{x}^{(t)})$      $\{\eta > 0 : \text{step size}\}$

**end for**



# Stochastic Gradient Descent

- ▶ Empirical risk minimization: additive objective

$$\text{minimize} \quad f(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x})$$

- ▶ Stochastic Gradient Descent (SGD)

initialize  $\mathbf{x}^{(0)} \in \mathbb{R}^m$

**for**  $t = 0:T-1$  **do**

sample  $i \stackrel{\text{uni}}{\sim} \{1 \dots n\}$

update  $\mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \eta \nabla f_i(\mathbf{x}^{(t)})$

**end for**

# Stochastic Gradient Descent

$$\text{SGD update } \mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \eta \nabla f_i(\mathbf{x}^{(t)})$$

- ▶ **Idea:** Cheap but **unbiased** estimate of the gradient
  - ▶  $\mathbf{E} \nabla f_i(\mathbf{x}) = \nabla f(\mathbf{x})$  over random choice of  $i$
  - ▶ downside: variance = randomness in descent directions
- ▶ computing  $\nabla f_i(\mathbf{x})$  is a factor  $n$  cheaper than computing  $\nabla f(\mathbf{x})$
- ▶ convergence: decreasing stepsize  $\eta \propto \frac{1}{t}$



## Section 3

# Constrained Optimization

# Constrained Optimization

Constrained Optimization Problem

$$\begin{array}{ll} \text{minimize} & f(\mathbf{x}) \\ \text{subject to} & \mathbf{x} \in Q \end{array}$$




# Projected Gradient Descent

Idea: project onto  $Q$  after each step:

$$P_Q(\mathbf{x}) := \arg \min_{\hat{\mathbf{x}} \in Q} \|\hat{\mathbf{x}} - \mathbf{x}\|$$

Projected gradient update

$$\mathbf{x}^{(t+1)} \leftarrow P_Q \left[ \mathbf{x}^{(t)} - \eta \nabla f(\mathbf{x}^{(t)}) \right]$$



# Lagrangian Function

## Primal optimization problem

$$\begin{array}{ll}\text{minimize} & f(\mathbf{x}) \\ \text{subject to} & g_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, p \\ & h_j(\mathbf{x}) = 0, \quad j = 1, \dots, q\end{array}$$

## Lagrangian

$$L(\mathbf{x}, \lambda, \nu) := f(\mathbf{x}) + \sum_{i=1}^p \lambda_i g_i(\mathbf{x}) + \sum_{j=1}^q \nu_j h_j(\mathbf{x})$$

- $\lambda_i \geq 0, \nu_k \in \mathbb{R}$ : Lagrange multipliers



# Lagrangian Dual

- ▶ Lagrange **dual function**:

$$\mathcal{D}(\lambda, \nu) := \inf_{\mathbf{x}} L(\mathbf{x}, \lambda, \nu) \quad \in \mathbb{R}$$

- ▶ for any feasible  $\mathbf{x}$ :  $\nu_i h_i(\mathbf{x}) = 0$  and  $\lambda_j g_j(\mathbf{x}) \leq 0$
- ▶ hence:  $\mathcal{D}(\lambda, \nu) \leq f(\mathbf{x})$ ,  $\mathbf{x}$ : feasible
- ▶ Lagrange **dual problem**: best lower bound on  $f(\mathbf{x}^*)$

$$(\lambda^*, \nu^*) = \arg \max_{\lambda \geq 0, \nu} \mathcal{D}(\lambda, \nu)$$

$$\mathcal{D}(\lambda, \nu) \leq \mathcal{D}(\lambda^*, \nu^*) \leq f(\mathbf{x}^*) \leq f(\mathbf{x})$$

# Karush-Kuhn-Tucker Conditions

- ▶ assume  $\mathbf{x}^*$  is a local minimum,  $f, g_i, h_j$  continuously diff at  $\mathbf{x}^*$
- ▶ ... under some regularity conditions (e.g. affine  $g_i$  and  $h_j$ )
- ▶ **KKT** conditions hold for  $\lambda^*$  and  $\nu^*$

$$(1) \quad \nabla_{\mathbf{x}} L(\mathbf{x}^*; \lambda^*, \nu^*) = \mathbf{0} \quad \text{1st order optimality}$$

$$(2) \quad g_i(\mathbf{x}^*) \leq 0, \quad h_j(\mathbf{x}^*) = 0, \quad \lambda_i^* \geq 0, \quad \text{feasability}$$

$$(3) \quad \lambda_i^* g_i(\mathbf{x}^*) = 0 \quad \text{complementary slackness}$$

- ▶ it is implied that  $f(\mathbf{x}^*) = L(\mathbf{x}^*; \lambda^*, \nu^*)$

# Strong Duality

- Duality gap:

$$\Delta = \min_{\mathbf{x}} f(\mathbf{x}) - \max_{\lambda \geq \mathbf{0}, \nu} \mathcal{D}(\lambda, \nu) \geq 0$$

- $\Delta = 0$ , if primal problem is **convex** and under additional conditions
- if  $\Delta = 0$ , then

$$\mathcal{D}(\lambda^*, \nu^*) = \min_{\mathbf{x}} L(\mathbf{x}, \lambda^*, \nu^*) = L(\mathbf{x}^*, \lambda^*, \nu^*) = f(\mathbf{x}^*)$$

- can recover  $\mathbf{x}^*$  by unconstrained minimization of  $L(\mathbf{x}, \lambda^*, \nu^*)$
- dual problem has very simple constraints!



## Section 4

# Convex Optimization

# Convex Set

A set  $Q$  is *convex* if the line segment between any two points of  $Q$  lies in  $Q$ , i.e., if for any  $\mathbf{x}, \mathbf{y} \in Q$  and any  $\theta$  with  $0 \leq \theta \leq 1$ , we have

$$\theta \mathbf{x} + (1 - \theta) \mathbf{y} \in Q.$$



\*Figure 2.2 from S. Boyd, L. Vandenberghe

**Left** Convex.

**Middle** Not convex, since line segment not in set.

**Right** Not convex, since some, but not all boundary points are contained in the set.

# Properties of Convex Sets

- ▶ Intersections of convex sets are convex
- ▶ Projections onto convex sets are *unique*.  
(and often efficient to compute)

$$\text{recall } P_Q(\mathbf{x}') := \arg \min_{\mathbf{y} \in Q} \|\mathbf{y} - \mathbf{x}'\|$$

# Convex Function


**Epigraph:** the *graph* of a function  $f : \mathbb{R}^m \rightarrow \mathbb{R}$  is defined as

$$\{(\mathbf{x}, f(\mathbf{x})) \mid \mathbf{x} \in \text{dom} f\},$$

The *epigraph* of a function  $f : \mathbb{R}^m \rightarrow \mathbb{R}$  is defined as

$$\{(\mathbf{x}, t) \mid \mathbf{x} \in \text{dom} f, f(\mathbf{x}) \leq t\},$$

A function is convex *iff* its epigraph is a convex set.



Convex?

\*Figure 3.5 from S. Boyd, L. Vandenberghe

# Convex Function

Examples of convex functions

- ▶ Linear functions:  $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x}$
- ▶ Affine functions:  $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} + b$
- ▶ Exponential:  $f(\mathbf{x}) = e^{\alpha \mathbf{x}}$
- ▶ Norms. Every norm on  $\mathbb{R}^m$  is convex.

Convexity of a norm  $f(\mathbf{x})$

- ▶ triangle inequality  $f(\mathbf{x} + \mathbf{y}) \leq f(\mathbf{x}) + f(\mathbf{y})$
- ▶ homogeneity of a norm  $f(a\mathbf{x}) = |a|f(\mathbf{x})$

$$f(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) \leq f(\theta \mathbf{x}) + f((1 - \theta) \mathbf{y}) = \theta f(\mathbf{x}) + (1 - \theta) f(\mathbf{y}).$$



# Convex Optimization

- ▶ **Convex Optimization** problems

$$\min f(\mathbf{x}) \quad \text{s.t.} \quad \mathbf{x} \in Q$$

where both

- ▶  $f$  is a convex function
  - ▶  $Q$  is a convex set (note:  $\mathbb{R}^m$  is convex)
- 
- ▶ Properties of Convex Optimization Problems
    - ▶ **every local minimum is a global minimum**

# Solving Convex Optimization Problems (provably)

- ▶ For convex optimization problems, all algorithms
  - ▶ Coordinate Descent
  - ▶ Gradient Descent
  - ▶ Stochastic Gradient Descent
  - ▶ Projected Gradient Descent (projections onto convex sets do work!)

do **converge** to the global optimum! (assuming  $f$  differentiable)

- ▶ **Theorem:** For convex problems, the **convergence rate** of the above four algorithms is (at least) proportional to  $\frac{1}{t}$ , i.e.


$$f(\mathbf{x}^{(t)}) - f(\mathbf{x}^*) \leq \frac{c}{t}$$

caveat: SGD rate can be  $1/\sqrt{t}$  if  $f$  is not strongly convex

# Sub-Gradient Descent

- ▶ What if  $f$  is not differentiable?
- ▶ Work with Sub-gradient:  $\mathbf{g} \in \mathbb{R}^m$  is a **subgradient** of  $f$  at  $\mathbf{x}$  if

$$f(\mathbf{y}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{y} - \mathbf{x}) \quad \text{for all } \mathbf{y}$$



# Sub-Gradient Descent

- ▶ **Subgradient Descent:** in algorithms, replace the gradient with a subgradient.
- ▶ **Theorem:** For convex problems, the convergence rate of [plain or projected] subgradient descent is (at least) proportional to  $\frac{1}{\sqrt{t}}$ , i.e.

$$f(\mathbf{x}^{(t)}) - f(\mathbf{x}^*) \leq \frac{c}{\sqrt{t}}$$

## Section 5

### Convex Relaxation

# Convex Relaxation

- ▶ For cases where  $Q$  is not convex ...
  - ▶ ... we can aim to find  $Q' \supset Q$
  - ▶ ... such that  $Q'$  is convex
  - ▶ ... and then solve the **convex relaxation**

$$\min f(\mathbf{x}) \quad \text{subject to} \quad \mathbf{x} \in Q'$$

- ▶ ... finally we can try to find a close(st) point in  $Q$

# Convex Relaxation for Matrix Completion

- ▶ Replace non-convex rank constraint by convex norm constraint.
- ▶ Nuclear norm

$$\|\mathbf{A}\|_* = \sum_i \sigma_i, \quad \sigma_i : \text{singular values}$$

- ▶ Exact reconstruction:

$$\min_{\mathbf{B}} \|\mathbf{B}\|_*, \quad \text{subject to} \quad A_{ij} = B_{ij} \quad \forall (i, j) \in \mathcal{I}$$

- ▶ Approximate reconstruction (unconstrained problem)

$$\min_{\mathbf{B}} \|\mathbf{B}\|_* + \lambda \sum_{(i,j) \in \mathcal{I}} (a_{ij} - b_{ij})^2$$

# Compressed Sensing for Matrix Completion

- ▶ **Theorem:** exact reconstruction of a  $\mathbb{R}^{m \times n}$  rank  $k$  matrix ( $m \leq n$ ) is possible with probability  $1 - n^{-3}$ , if it is strongly incoherent with parameter  $\mu$  (spread of singular values) and as long as

$$|\mathcal{I}| \geq C\mu^4 k^2 n (\log n)^2, \quad \text{typically} \quad \mu = \mathbf{O}(\sqrt{\log n})$$

- ▶ Candes, Tao: The power of convex relaxation: Near-optimal matrix completion, 2010
- ▶ explains, why  $\|\cdot\|_*$  minimization works well in practice!